



# Sharding and Scale-out using MySQL Fabric

Mats Kindahl ([mats.kindahl@oracle.com](mailto:mats.kindahl@oracle.com))  
Principal Senior Software Developer

# Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decision. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.



# Presentation Outline

- Introducing MySQL Fabric
- Architecture for High-Availability
- Connecting to a MySQL Fabric Farm
- Architecture for Sharding
- Summary and Closing Remarks

# MySQL Fabric

An extensible and easy-to-use framework for managing a farm of MySQL servers supporting high-availability and sharding



# What does all that mean?

- “Farm” Management System
  - **Farm:** Collection of components
  - Distributed Framework
- Framework
  - Procedure execution
  - State store
  - Transaction Routing
- Extensible
  - Extensions are first-class
  - High-Availability Groups
  - “Semi-Automatic” Sharding
- Written in Python
- Latest Release 1.4.2
  - Release Candidate
- Open Source
  - **You** can participate
  - Suggest features
  - Report bugs
  - Contribute patches
- MySQL 5.6 is focus

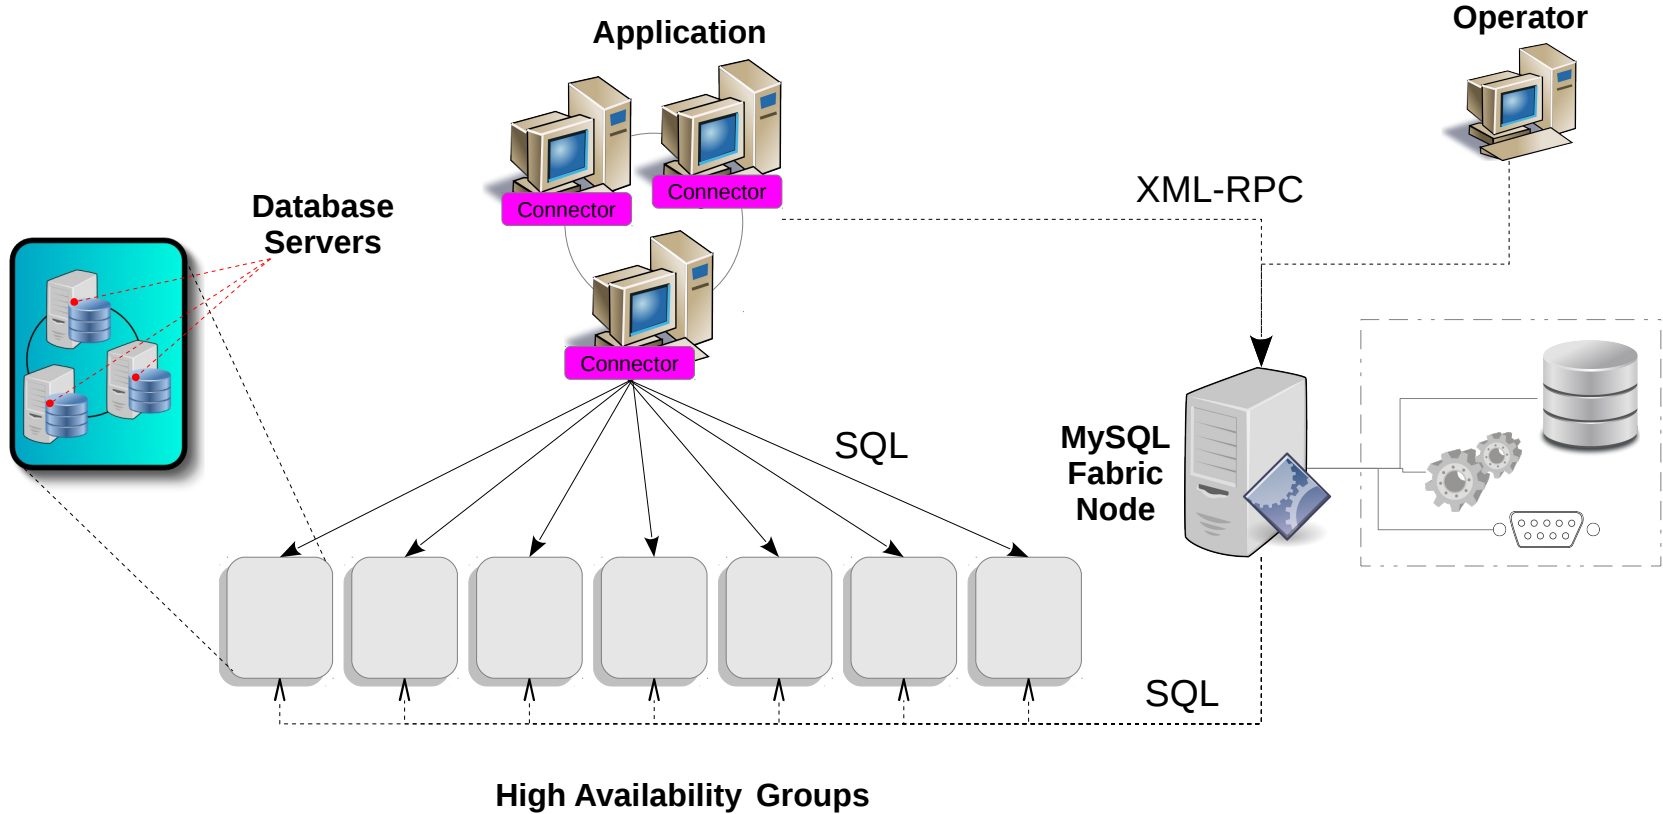
# MySQL Fabric: Goals & Features

- Decision logic in connector
  - Eliminate one network hop
  - Reducing network load
- Connector API Extensions
  - Support Transactions
  - Support full SQL
- Load Balancing
  - Read-Write Split
  - Round-robin
- Multi-Table Sharding
- Sharding Functions
  - Range
  - (Consistent) Hash
- Shard Operations
  - Shard move
  - Shard split
- Global Updates
  - Global tables
  - Schema updates

# A Brief History of MySQL Fabric

- MySQL Fabric 1.4.0
  - September, 2013
  - First public release
  - High-Availability Groups
  - Slave Promotion
  - Range and Hash Sharding
  - Shard move and split
  - Connector/Python Support
  - Connector/J Support
  - Connector/PHP Support
- MySQL Fabric 1.4.1
  - December, 2013
  - Alpha release
  - Sharding refactorings
- MySQL Fabric 1.4.2
  - April, 2014
  - Release Candidate
  - Distributed Failure Detection
  - Credentials
  - Weighted Round-Robin

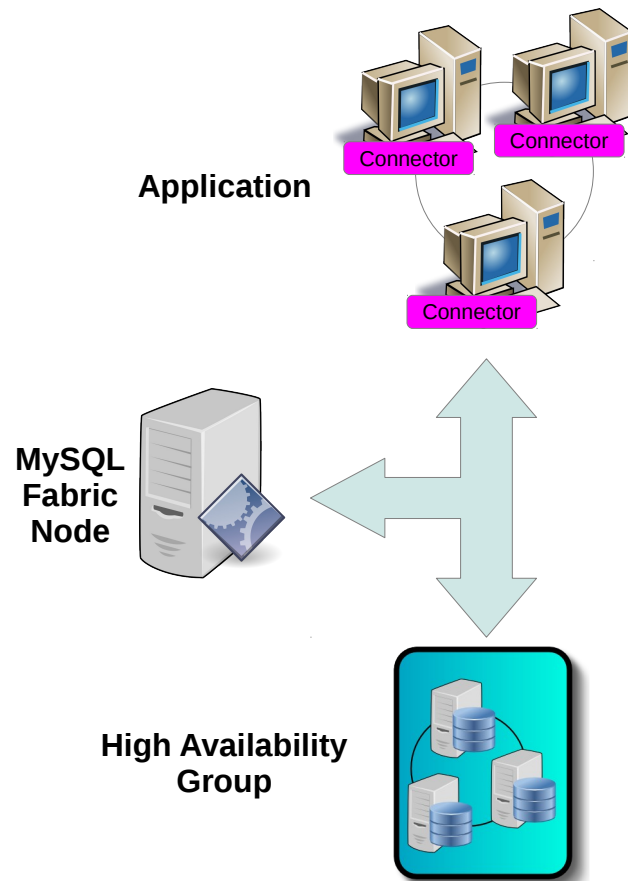
# Birds-eye View



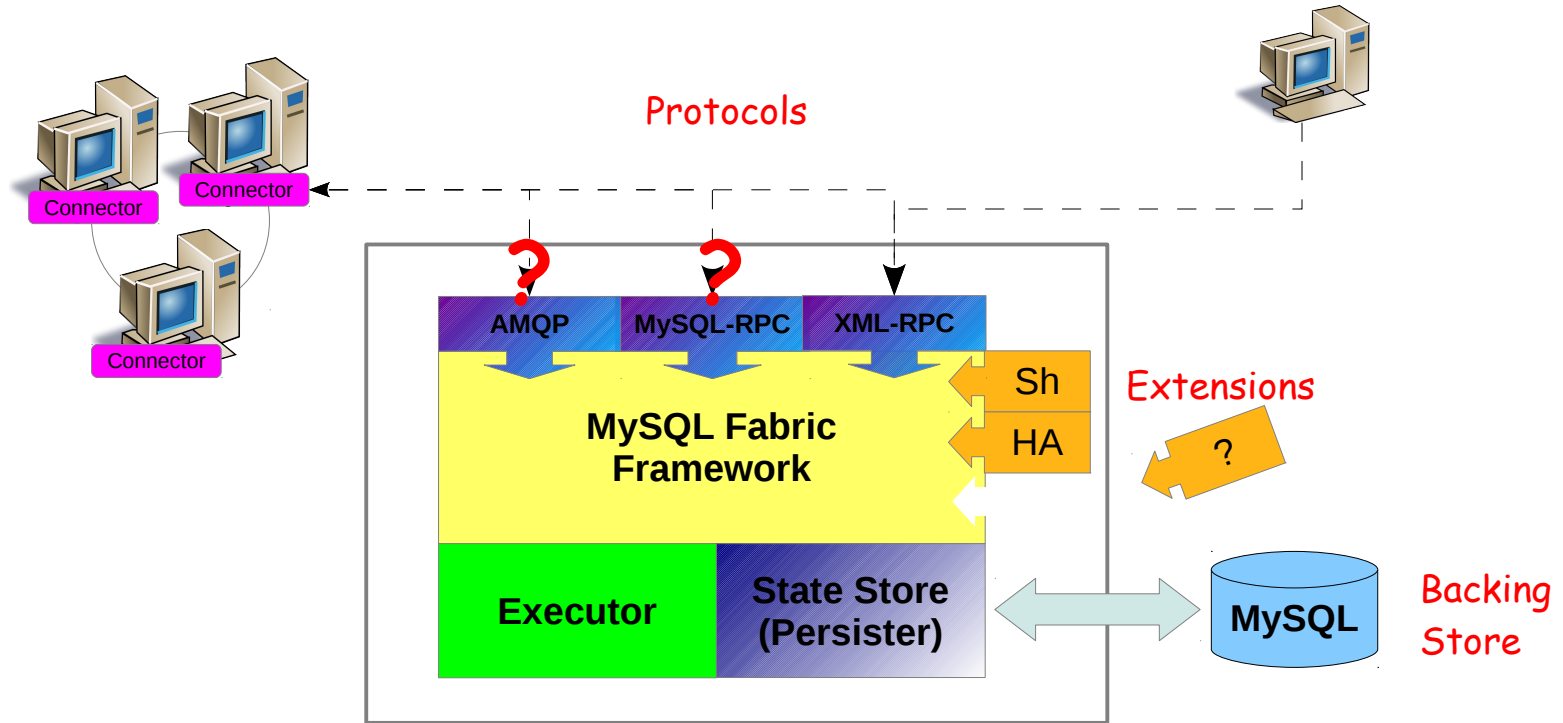


# High-Level Components

- Fabric-aware Connectors
  - Python, PHP, and Java
  - Enhanced Connector API
- MySQL Fabric Node
  - Manage information about farm
  - Provide status information
  - Execute procedures
- MySQL Servers
  - Organized in High-Availability Groups
  - Handling application data



# MySQL Fabric Node Architecture



# MySQL Fabric: Prerequisites

- MySQL Servers (version 5.6.10 or later)
  - Backing store database server
  - Application database servers
- Python 2.6 or 2.7
  - No support for 3.x yet
- MySQL Utilities 1.4
  - Available at <https://dev.mysql.com/downloads/tools/utilities>
  - “Development release” tab

# MySQL Fabric: Configuration

- Backing Store
  - MySQL server
  - Persistent storage for state
  - Storage engine-agnostic
- Protocol
  - Address where node will be
  - Currently only XML-RPC
- Logging
  - Chatty: **INFO** (default)
  - Moderate: **WARNING**
  - URL for rotating log

```
[storage]
address = localhost:3306
user = fabric
password =
database = fabric
```

```
[servers]
user = fabric
password =
```

```
[protocol.xmlrpc]
address = localhost:32274
threads = 5
disable_authentication = yes
```

```
[logging]
level = INFO
url = file:///var/log/fabric.log
```

# MySQL Fabric: Basic Commands and Help

- Command Structure

```
mysqlfabric group command ...
```

- Getting help

```
mysqlfabric help  
mysqlfabric help commands  
mysqlfabric help manage  
mysqlfabric help manage setup
```

- MySQL Utilities Documentation:
  - <http://dev.mysql.com/doc/mysql-utilities/1.4/en/index.html>
- MySQL Fabric Documentation:
  - <http://dev.mysql.com/doc/mysql-utilities/1.4/en/fabric.html>

# Setting up and Tearing down MySQL Fabric

- Create and populate the necessary tables in backing store

```
mysqlfabric manage setup
```

- Remove the tables from backing store

```
mysqlfabric manage teardown
```

- Connects to the database server in “storage” section
  - Ensure that you have the necessary users and privileges

# Starting and Stopping MySQL Fabric

- Start MySQL Fabric node in foreground – print log to terminal

```
mysqlfabric manage start
```

- Start MySQL Fabric node in background – print log to file

```
mysqlfabric manage start --daemonize
```

- Stop MySQL Fabric node

```
mysqlfabric manage stop
```


# Architecture for High-Availability





# High-Availability Concepts




- Redundancy
  - Duplicate critical components
- Monitoring
  - Detecting failing components
  - Monitor load
- Procedures
  - Activate replacements
  - Distribute load



# High-Availability Group Concept


- Group of servers
  - Hardware redundancy
  - Data redundancy
- Generic Concept
  - Implementation-independent
  - Self-managed or externally managed
- Different Types
  - Primary-Backup (Master-Slave) **Done!**
  - Shared or Replicated Storage
  - MySQL Cluster

**Examples Only  
Not Implemented**



# High-Availability Group Concept

- Abstract Concept
  - Set of servers
  - Server attributes
- Connector Attributes
  - Connection information
  - **Mode:** read-only, read-write, ...
  - **Weight:** distribute load
- Management Attributes
  - **Status:** state/role of the server



Status: Primary  
Mode: Read-Write  
Host: server-  
1.example.com

# Create Groups and add Servers

- Define a group

```
mysqlfabric group create my_group
```

- Add servers to group

```
mysqlfabric group add my_group server1.example.com  
mysqlfabric group add my_group server2.example.com
```

# Activate High-Availability Group

- Promote one server to be primary

```
mysqlfabric group promote my_group
```


- Tell built-in failure detector to monitor group

```
mysqlfabric group activate my_group
```

# Distributed Failure Detector

*New in MySQL Fabric 1.4.2*

- Connectors report errors
  - Report that an error was noticed
  - Failover based on statistics
  - **report\_error(server, source, error)**
- Report failure
  - A server is known to have failed
  - Failover occurs immediately
  - **report\_fault(server, source, error)**



# Update Only Operations

*New in MySQL Fabric 1.4.2*

- **Situation:**
  - Server promotion is done elsewhere
  - Real situation does not match content of state-store
- **Problem:**
  - Need to update state store to match real situation
  - Should not touch application servers
- **Solution:** Use update-only option

```
mysqlfabric group promote my_group --update_only \  
--slave_uuid=29bf3b2d-b5ac-11e3-a383-58946b051f64
```

# Credentials in MySQL Fabric

*New in MySQL Fabric 1.4.2*

- Credentials
  - Digest Authentication
  - SSL connection
- Digest Authentication
  - RFC 2617
  - **Server:** Realm
  - **Client:** User + Password
- SSL connection
  - Setup Like MySQL Server

```
[storage]
address = localhost:3306
user = fabric
password = xyzzy
database = fabric

[protocol.xmlrpc]
address = localhost:32274
threads = 5
disable_authentication = no
realm = MySQL Fabric
user = admin
password = xyzzy
ssl_ca = /etc/mysql/fabric_ca.pem
ssl_key = /etc/mysql/fabric_key.pem
ssl_cert = /etc/mysql/fabric_cert.pem
```



# Connecting to a MySQL Fabric Farm

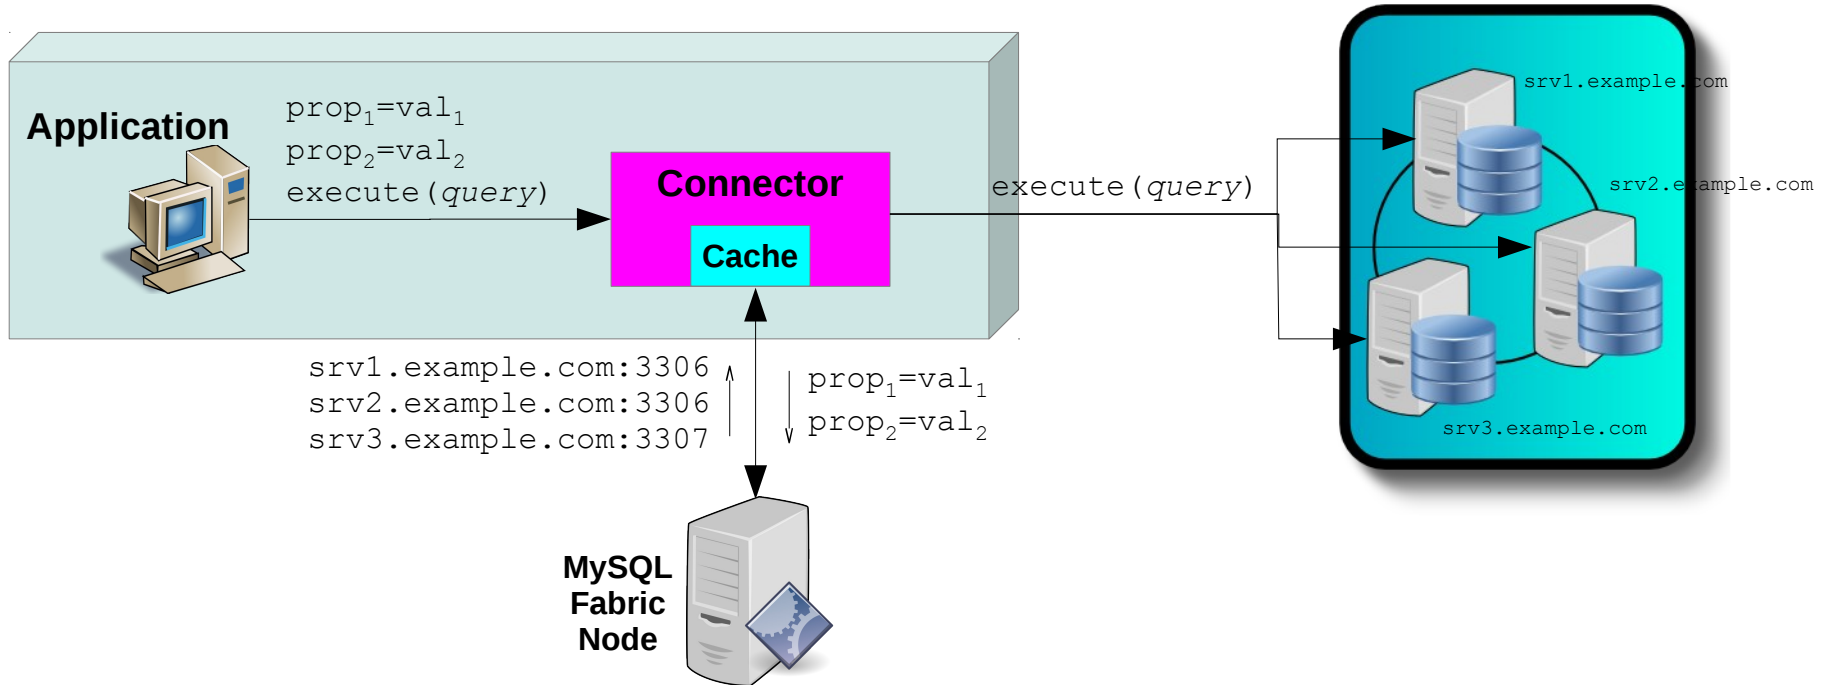




# Fabric-aware Connector

- Fabric-aware Connectors
  - Connector/J
  - Connector/Python
  - Connector/PHP
- Fabric-aware Frameworks
  - Doctrine
  - Hibernate
- In this presentation:
  - Connector/Python
- Connector API Extensions
  - Support Transactions
  - Support full SQL
- Decision logic in connector
  - Reducing network load
- Load Balancing
  - Read-Write Split
  - Distribute transactions

# Routing Transactions



# Fabric-aware Connector API

- Establish a “virtual” connection
  - Real server connection established lazily
- Provide connection information for the *Fabric node*
  - Connector will fetch information about servers

```
import mysql.connector

conn = mysql.connector.connect(
    fabric={"host": "fabric.example.com"},
    user='mats', password='xyzyz', database="employees"
)
```

# Enable Connector/Python Error Reporting

*New in Connector/Python 1.2.1*

- Connectors can report errors to Fabric node
  - Enable using `report_error`
  - Defaults to **False**
  - Require MySQL Fabric 1.4.2

```
import mysql.connector

conn = mysql.connector.connect(
    fabric={"host": "fabric.example.com"},
    user='mats', password='xyzzzy', database="employees",
    report_error=True,
)
```

# Connector API: Executing a Transaction

- Provide group name
  - **Property:** group
  - Fabric will compute candidate servers
- Provide transaction mode
  - **Property:** mode
  - Fabric will pick server in right mode

```
conn.set_property(group='my_group', mode=MODE_READWRITE)
cur = conn.cursor()
cur.execute("INSERT INTO employees VALUES (%s,%s,%s)",
            (emp_no, first_name, last_name))
cur.execute("INSERT INTO titles(emp_no,title,from_date)"
            " VALUES (%s,%s,CURDATE())",
            (emp_no, 'Intern'));
conn.commit()
```

Same as  
before

# Executing a Transaction

Hmm... looks like  
a read transaction

Where's the sharding key?

Ah, there it is!

Session state?

```
START TRANSACTION;  
SELECT salary INTO @s FROM salaries WHERE emp_no = 20101;  
SET @s = 1.1 * @s;  
INSERT INTO salaries VALUES (20101, @s);  
COMMIT;  
BEGIN;  
CALL update_salary(20202, @s);  
COMMIT;
```

What does this  
procedure update?

Oops... it was a  
write transaction!

Transaction done!

Clear session state?

What about connection pools?  
Application error?

New transaction! Different connection?  
What about the session state?


# Architecture for Sharding





# Benefits of Sharding

- Write scalability
  - Can handle more writes
- Large data set
  - Database too large
  - Does not fit on single server
- Improved performance
  - Smaller index size
  - Smaller working set



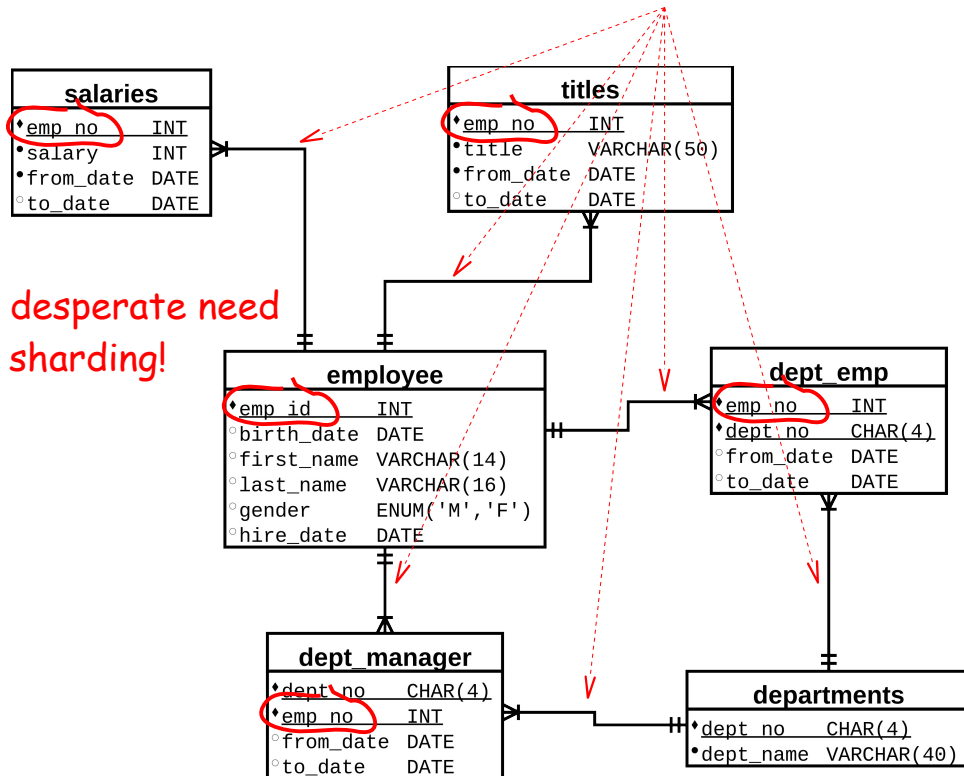
# MySQL Fabric: Sharding Goals & Features

- Connector API Extensions
  - Support Transactions
  - Support full SQL
- Decision logic in connector
  - Reducing network load
- Shard Multiple Tables
  - Using same key
- Global Updates
  - Global tables
  - Schema updates
- Sharding Functions
  - Range
  - (Consistent) Hash
- Shard Operations
  - Using built-in executor
  - Shard move
  - Shard split

# Sharded Tables

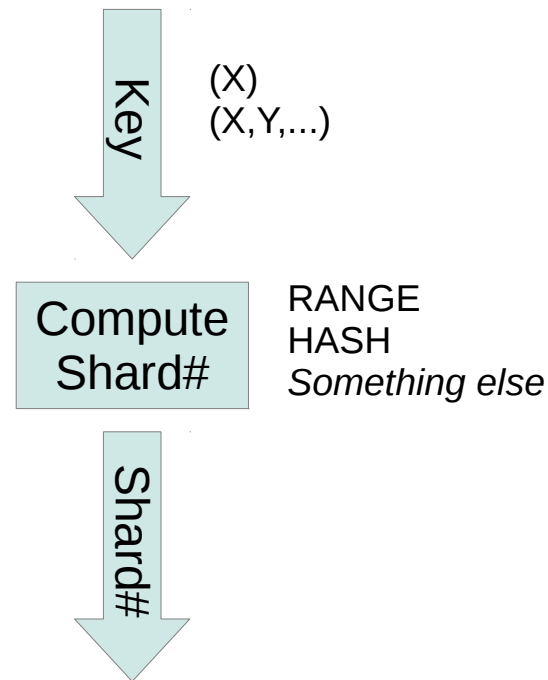
Table	Rows
salaries	284 404 700
titles	44 330 800
employees	30 002 400
dept_emp	33 160 300
dept_manager	2 400
departments	900

In desperate need  
of sharding!

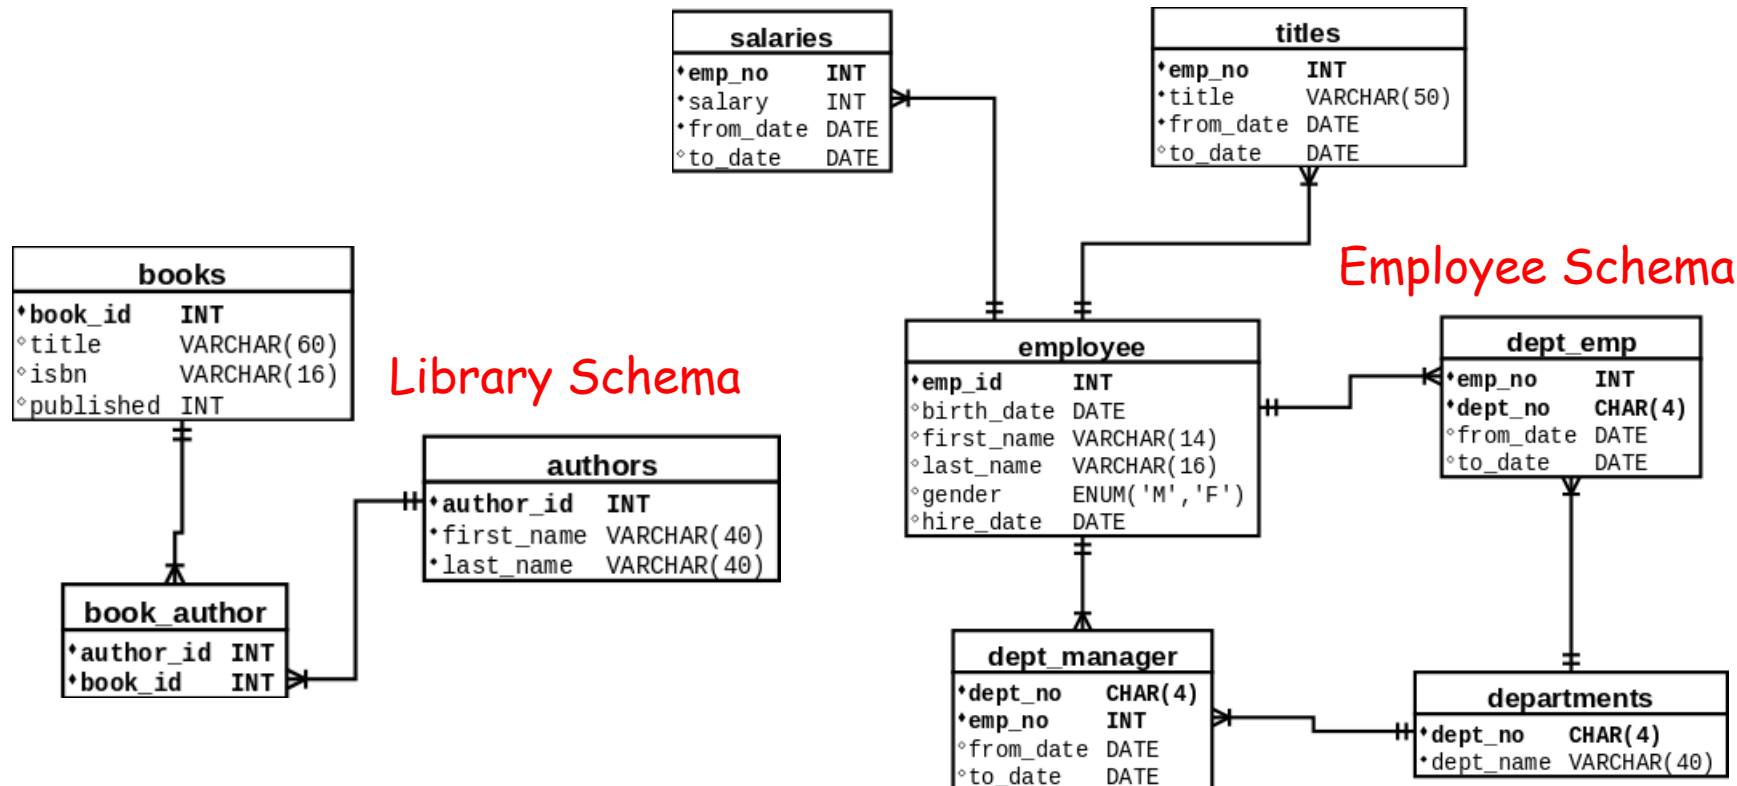


# Mapping the Sharding Key

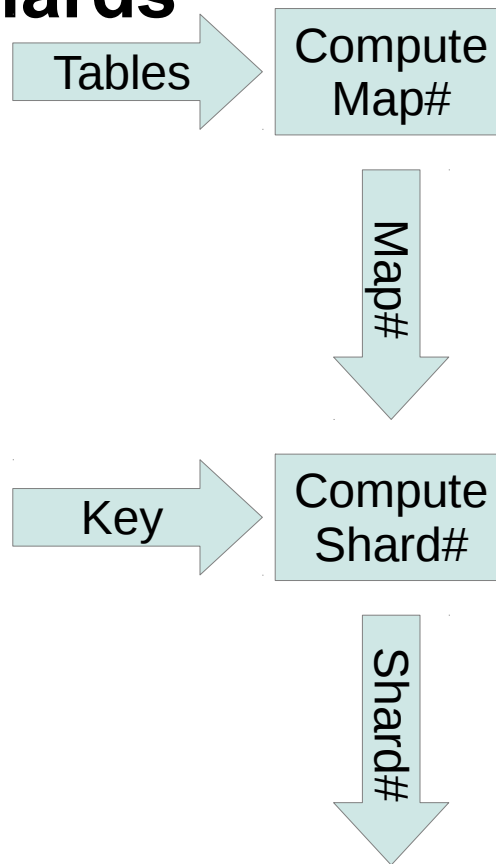
- What is a sharding key?
  - Single column
  - Multi column
    - Same table?
    - Different tables?
- How is the key transformed?
  - Hash
  - Range
  - User-defined



# Sharded Tables: Multiple Mappings



# Digression: Computing Shards



- Multiple Mappings
  - Which mapping to use?
  - Application don't care
    - ... but know tables in transaction
  - Currently only one mapping
- Computing shard requires
  - Tables + sharding key
  - Map# + sharding key
- Enhanced Connector API

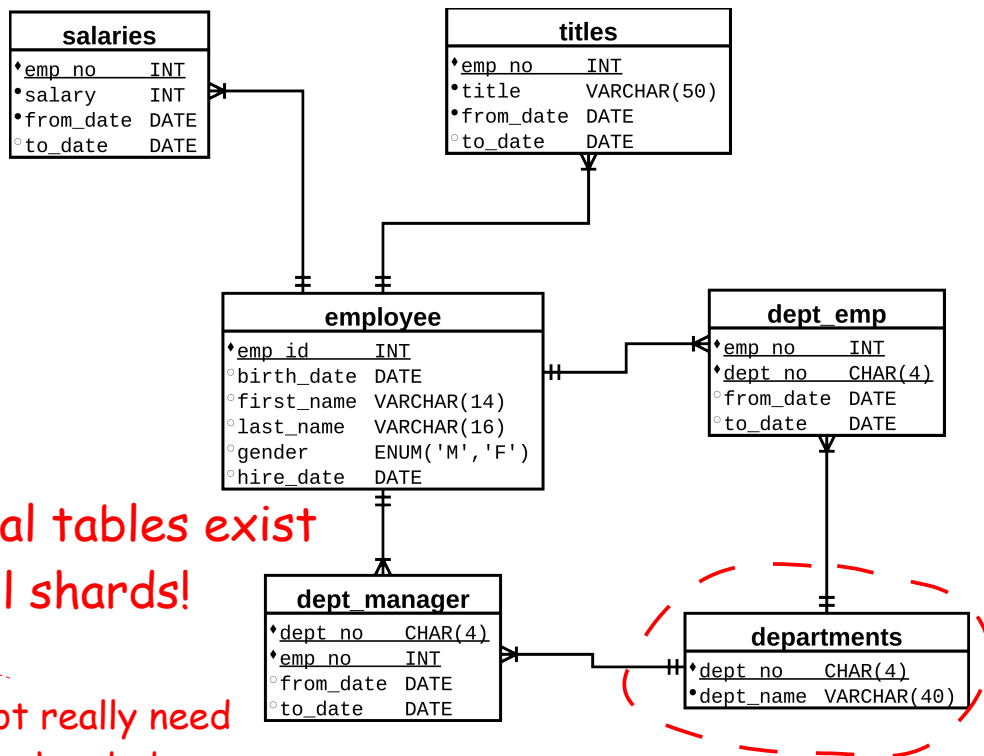
# Multi-table Query with Sharded Tables

```
SELECT first_name, last_name, salary
FROM salaries JOIN employees USING (emp_no)
WHERE emp_no = 21012
      AND CURRENT_DATE BETWEEN from_date AND to_date;
```

- Referential Integrity Constraint
  - Example query joining salaries and employees
  - Same key, same shard: co-locate rows for same user
- JOIN normally based on equality
  - Using non-equality defeats purpose of foreign key

# Global Tables

Table	Rows
salaries	284 404 700
titles	44 330 800
employees	30 002 400
dept_emp	33 160 300
dept_manager	2 400
departments	900



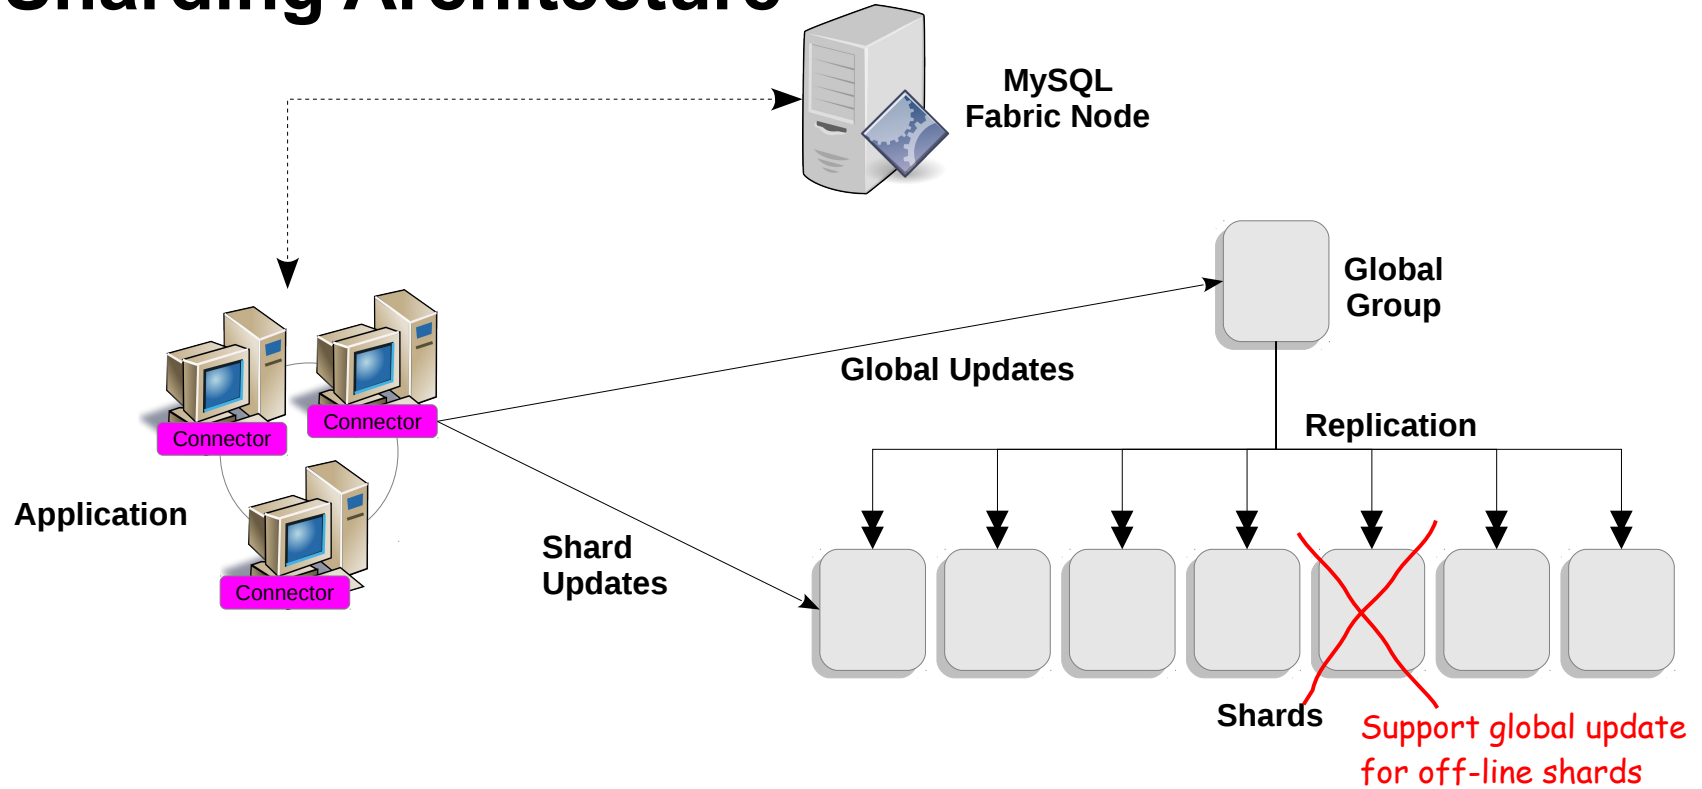


# Multi-table Query with Global Tables

```
SELECT first_name, last_name, GROUP_CONCAT(dept_name)
FROM employees JOIN dept_emp USING (emp_no)
              JOIN departments USING (dept_no)
WHERE emp_no = 21012 GROUP BY emp_no;
```

- JOIN with **departments** table
  - Has no employee number, hence no sharding key
  - Table need to be present on all shards
- But... how do we update global tables?

# Sharding Architecture



# MySQL Fabric: Sharding Setup

- Set up some groups
  - `my_global` – for global updates
  - `my_group.*` – for the shards
  - Add servers to the groups
- Create a shard mapping
  - A “distributed database”
  - Give information on what tables are sharded
- Add shards
  - Mapping keys to shards

# MySQL Fabric: Set up Shard Mapping

- Define shard mapping

```
mysqlfabric sharding \  
    create_definition hash my_global
```

Will return a  
shard map identifier

- Add tables that should be sharded

```
mysqlfabric sharding add_table 1 \  
    employees.employees emp_no  
mysqlfabric sharding add_table 1 \  
    employees.salaries emp_no
```

Shard map identifier

- *Tables not added are considered global*

# MySQL Fabric: Add Shards

- Add shards to shard mapping

```
mysqlfabric sharding add_shard 1 \  
    "my_group.1,...,my_group.N" --state=ENABLED
```

Shard map identifier



# MySQL Fabric: Moving and Splitting Shards

- Moving a shard from one group to another

```
mysqlfabric sharding move 5 my_group.5
```

- Splitting a shard into two pieces (hash)

```
mysqlfabric sharding split 5 my_group.6
```

Shard ID



# Connector API: Shard Specific Query

- Provide tables in query
  - **Property:** tables
  - Fabric will compute map
- Provide sharding key
  - **Property:** key
  - Fabric will compute shard

```
conn.set_property(tables=['employees.employees', 'employees.titles'],
                  key=emp_no)
cur = conn.cursor()
cur.execute("INSERT INTO employees VALUES (%s,%s,%s)",
            (emp_no, first_name, last_name))
cur.execute("INSERT INTO titles(emp_no, title, from_date)"
            " VALUES (%s, %s, CURDATE())",
            (emp_no, 'Intern'));
conn.commit()
```

# Connector API: Shard Specific Query

- Provide tables in query
  - **Property:** tables
  - Fabric will compute map
- Provide sharding key
  - **Property:** key
  - Fabric will compute shard

```
conn.set_property(tables=['employees.employees', 'employees.titles'],
                  key=emp_no)
cur = conn.cursor()
cur.execute(
    "SELECT first_name, last_name, title"
    "  FROM employees JOIN titles USING (emp_no)"
    " WHERE emp_no = %d", (emp_no,))
for row in cur:
    print row[0], row[1], ",", row[2]
```

Join queries are sent to correct  
shard and executed there



# Connector API: Global Update

- Provide tables in query
  - **Property:** tables
  - Fabric will compute map
  - (Not necessary)
- Set global scope
  - **Property:** scope
  - Query goes to global group

```
conn.set_property(tables=['employees.titles'], scope='GLOBAL')  
cur = conn.cursor()  
cur.execute("ALTER TABLE employees.titles ADD nickname VARCHAR(64)")
```

# Closing Remarks



# What do we have now?

- MySQL Farm Management
  - High-Availability
  - Sharding
- High-Availability
  - Group Concept
  - Slave promotion
- Sharding
  - Range and hash sharding
  - Shard move and shard split
- Connector APIs
  - Transaction properties
  - “Virtual” connections
- Enhanced Connectors
  - Connector/Python
  - Connector/PHP
  - Connector/J
- Command-line Interface
- XML-RPC Interfaces
- Distributed failure detector
  - Connectors report failures
  - Custom failure detectors
- Credentials
  - RFC 2617
  - SSL support

# Thoughts for the Future

- Connector multi-cast
  - Scatter-gather
  - UNION of result sets
  - More complex operations?
- Extension interfaces
  - Improve extension support
  - Improve procedures support
- Command-line interface
  - Improving usability
  - Focus on ease-of-use
- More protocols
  - MySQL-RPC Protocol?
- More frameworks?
- More connectors?
  - C/C++?
  - Fabric-unaware connectors?
- More HA group types
  - DRBD
  - MySQL Cluster

# Thoughts for the Future

- “Transparent” Sharding
  - Single-query transactions?
  - Speculative execution?
  - Cross-shard join?
- Multiple shard mappings
  - Independent tables
- Multi-way shard split
  - Efficient initial sharding
  - Better use of resources
- High-availability executor
  - Node failure stop execution
  - Replicated State Machine
    - Paxos?
    - Raft?
  - Continue execution on other Fabric node
- Session Consistency
  - We have a distributed database
  - It should look like a single database

# Reading for the Interested

- MySQL Forum: *Fabric, Sharding, HA, Utilities*  
<http://forums.mysql.com/list.php?144>
- MySQL Fabric Documentation  
<http://dev.mysql.com/doc/mysql-utilities/1.4/en/fabric.html>
- Migrating From an Unsharded to a Sharded Setup  
<http://vnwrites.blogspot.com/2013/09/mysqlfabric-sharding-migration.html>
- Configuring and running MySQL Fabric  
<http://alfranio-distributed.blogspot.com/2014/03/mysqlfabric-installation.html>

# Want to contribute?

- Check it  
... and send us use-case and feature suggestions
- Test it  
... and send comments to the forum
- Break it  
... and send in bugs to <http://bugs.mysql.com>

# Keeping in Touch

Mats Kindahl

**Twitter:** @mkindahl

<http://mysqlmusings.blogspot.com>

Alfranio Correia

**Twitter:** @alfranio

<http://alfranio-distributed.blogspot.com>

Narayanan Venkateswaran

**Twitter:** @vn\_tweets

<http://vnwrites.blogspot.com>

Geert Vanderkelen

**Twitter:** @geertjanvdk

<http://geert.vanderkelen.org>



# MySQL Connect @ Oracle

## OpenWorld

### Sept 29 – Oct 2



**Call for Proposals Open  
Until April 15!**



#### Special Offer:

**Register Before  
May 2 and Save an  
Extra US\$100!  
Code: DPER**

- Keynotes
- Conferences Sessions
- Birds-of-a-feather sessions
- Tutorials
- Hands-on Labs
- Demos
- Receptions
- OpenWorld Extensive Content

# Thank you!

